

Programming with weak synchronization models

Guest lecture ID2203

Peter Van Roy
Christopher Meiklejohn
Annette Bieniusa



Overview of the lesson

- Motivation and principles
 - “As easy as strong consistency, as efficient as weak consistency”
 - A sweet spot: Strong Eventual Consistency
- Convergent data structures
 - Conflict-free replicated data types (CRDTs)
- Lasp
 - Programming language and platform based on composing CRDTs
- Antidote
 - Causal transactional database based on CRDTs

Guest lecturers

- This lesson is brought to you by:

- Peter Van Roy, Université catholique de Louvain
- Christopher Meiklejohn, Université catholique de Louvain
- Annette Bieniusa, Technische Universität Kaiserslautern



- This research is being done in two European projects:



EU FP7
2013-2016



EU H2020
2017-2019



Both easy and efficient

- One of the holy grails of distributed systems is to make them both **easy to program** and **efficient to execute**
- **Strong consistency** (linearizability) is easy to program but inefficient
- **Eventual consistency** (operations eventually complete) is efficient to execute but hard to program
- Can we get the best of both worlds?
 - **Synchronization-free programming** aims to combine the ease of strong consistency with the efficiency of eventual consistency
 - How can this work?

Back to basics

- Distributed system = a collection of networked computing nodes that behaves like **one system** (= **consistency model**)
- To make this work, the nodes will coordinate with each other according to **well-defined rules** (= **synchronization algorithm**)
- For example, a **reliable broadcast algorithm** guarantees the **all-or-none property**: all correct nodes deliver, or none do

How far can we go?

- We would like the **consistency model to be as strong as possible** (easy to program) and the **synchronization algorithm to be as weak as possible** (efficient to execute)
- Let's try the extreme case: the weakest possible synchronization is no synchronization (no rules), which enforces no consistency at all!
 - So it's clear we need *some* synchronization
 - How little can we get away with?

A sweet spot: SEC

- **Strong Eventual Consistency (SEC)**
 - The data structure is defined so that n replicas that receive the same updates (in any order) have equivalent state
 - Synchronization is **eventual replica-to-replica communication**
- This consistency model is surprisingly powerful
 - It supports a programming model that resembles a concurrent form of functional programming
 - It handles both nondeterminism and nonmonotonicity
 - It has an efficient, resilient implementation

Let's exploit SEC!

- In the rest of the lesson we will see how far we can go with Strong Eventual Consistency
 - Convergent data structures (CRDTs)
 - Programming by composing CRDTs (Lasp)
 - Causally consistent transactions on CRDTs (Antidote)
 - Applications